

1 Useful concepts and Introduction

1.1 Usefull math

φ is convex $\Rightarrow \varphi(\mathbb{E}[X]) \leq \mathbb{E}[\varphi(X)]$

Hoeffding: $Z_1, \dots, iid, Z_i \in [0, C], \mathbb{E}[Z_i] = \mu$
 $\Rightarrow P\left(\left|\mu - \frac{1}{n} \sum_{i=1}^n Z_i\right| > \epsilon\right) \leq 2 \exp(-2n \frac{\epsilon^2}{C}) \leq \delta$
 $\Rightarrow n \geq \frac{C}{2\epsilon^2} \log \frac{2}{\delta}$

Robbins Monro $\alpha_t \xrightarrow{RM} 0: \sum \alpha_t = \infty, \sum \alpha_t^2 < \infty$

1.2 Multivariate Gaussian

$$f(x) = \frac{1}{2\pi\sqrt{|\Sigma|}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1}(x-\mu)}$$

Suppose we have a Gaussian random vector

$$\begin{bmatrix} X_A \\ X_B \end{bmatrix} \sim \mathcal{N}\left(\begin{bmatrix} \mu_A \\ \mu_B \end{bmatrix}, \begin{bmatrix} \Sigma_{AA} & \Sigma_{AB} \\ \Sigma_{BA} & \Sigma_{BB} \end{bmatrix}\right) \Rightarrow X_A|X_B = x_B \sim \mathcal{N}(\mu_A + \Sigma_{AB}\Sigma_{BB}^{-1}(x_B - \mu_B), \Sigma_{AA} - \Sigma_{AB}\Sigma_{BB}^{-1}\Sigma_{BA})$$

1.3 Information Theory elements:

Entropy: $H(X) \doteq -\mathbb{E}_{x \sim p_X} [\log p_X(x)]$

$$H(X|Y) \doteq -\mathbb{E}_{(x,y) \sim p_{(X,Y)}} [\log p_{Y|X}(y|x)]$$

$$\text{if } X \sim \mathcal{N}(\mu, \Sigma) \Rightarrow H(X) = \frac{1}{2} \log[(2\pi e)^d \det(\Sigma)]$$

Chain Rule: $H(X, Y) = H(Y|X) + H(X)$

Mutual Info: $I(X, Y) \doteq KL(p_{(X,Y)} \| p_X p_Y)$

$$I(X, Y) = H(X) - H(X|Y)$$

if $X \sim \mathcal{N}(\mu, \Sigma), Y = X + \epsilon, \epsilon \sim \mathcal{N}(0, \sigma^2 I)$:

$$\text{then } I(X, Y) = \frac{1}{2} \log \left[\det \left(I + \frac{1}{\sigma^2} \Sigma \right) \right]$$

1.4 Kullback-Leiber divergence

$$KL(p||q) = \mathbb{E}_p \left[\log \frac{p(x)}{q(x)} \right]$$

if $p_0 \sim \mathcal{N}(\mu_0, \Sigma_0), p_1 \sim \mathcal{N}(\mu_1, \Sigma_1) \Rightarrow KL(p_0||p_1)$

$$= \frac{1}{2} \left(\text{tr}(\Sigma_1^{-1} \Sigma_0) + (\mu_1 - \mu_0)^T \Sigma_1^{-1} (\mu_1 - \mu_0) - k + \log \frac{|\Sigma_1|}{|\Sigma_0|} \right)$$

$$\hat{q} = \arg \min_q KL(p||q) \Rightarrow \text{overconservative}$$

$$\hat{q} = \arg \min_q KL(q||p) \Rightarrow \text{overconfident}$$

2 Bayesian Regression

$$w \sim N(0, \sigma_p^2 I), \epsilon \sim N(0, \sigma_n^2 I), y = Xw + \epsilon$$

$$y|w \sim N(Xw, \sigma_n^2 I)$$

$$w|y \sim N((X^T X + \lambda I)^{-1} X^T y, (X^T X + \lambda I)^{-1} \sigma_n^2)$$

3 Kalman Filter

$$\begin{cases} X_{t+1} = F X_t + \epsilon_t & \epsilon_t \sim N(0, \Sigma_x) \\ Y_t = H X_t + \eta_t & \eta_t \sim N(0, \Sigma_y) \end{cases} \quad X_1 \sim N(\mu_p, \Sigma_p)$$

Then if X_0 is Gaussian then $X_t|Y_{1:t} \sim N(\mu_t, \Sigma_t)$:

$$\mu_{t+1} = F \mu_t + K_{t+1}(y_{t+1} - H F \mu_t)$$

$$\Sigma_{t+1} = (I - K_{t+1} H)(F \Sigma_t F^T + \Sigma_x)$$

$$K_{t+1} = (F \Sigma_t F^T + \Sigma_x) H^T (H(F \Sigma_t F^T + \Sigma_x) H^T + \Sigma_y)^{-1}$$

4 Gaussian Processes

$$f \sim GP(\mu, k) \Rightarrow \forall \{x_1, \dots, x_n\} \forall n < \infty$$

$$[f(x_1) \dots f(x_n)] \sim N([\mu(x_1) \dots \mu(x_n)], K)$$

where $K_{ij} = k(x_i, x_j)$

2.1 Gaussian Process Regression

$f \sim GP(\mu, k)$ then: $f|y_{1:n}, x_{1:n} \sim GP(\tilde{\mu}, \tilde{k})$

$$\tilde{\mu}(x) = \mu(x) + K_{A,x}^T (K_{AA} + \epsilon I_n)^{-1} (y_A - \mu_A)$$

$$\tilde{k}(x, x') = k(x, x') - K_{A,x}^T (K_{AA} + \epsilon I_n)^{-1} K_{A,x'}$$

Where: $K_{A,x} = [k(x_1, x) \dots k(x_n, x)]^T$

$$[K_{AA}]_{ij} = k(x_i, x_j) \text{ and } \mu_A = [\mu(x_1) \dots \mu(x_n)]^T$$

4.2 Kernels

$k(x, y)$ is a kernel if it's symmetric semidefinite positive:

$\forall \{x_1, \dots, x_n\}$ then for the Gram Matrix

$$[K]_{ij} = k(x_i, x_j) \text{ holds } c^T K c \geq 0 \forall c$$

Some Kernels: (h is the bandwidth hyperp.)

$$\text{Gaussian (rbf): } k(x, y) = \exp\left(-\frac{\|x-y\|^2}{h^2}\right)$$

$$\text{Exponential: } k(x, y) = \exp\left(-\frac{\|x-y\|}{h}\right)$$

$$\text{Linear kernel: } k(x, y) = x^T y \text{ (here } K_{AA} = X X^T)$$

4.3 Optimization of Kernel Parameters

Given a dataset A , a kernel function $k(x, y; \theta)$.

$y \sim N(0, K_y(\theta))$ where $K_y(\theta) = K_{AA}(\theta) + \sigma_n^2 I$

$$\hat{\theta} = \arg \max_{\theta} \log p(y|X; \theta)$$

$$\text{In GP: } \hat{\theta} = \arg \min_{\theta} y^T K_y^{-1}(\theta) y + \log |K_y(\theta)|$$

We can from here $\nabla \downarrow$:

$$\nabla_{\theta} \log p(y|X; \theta) = \frac{1}{2} \text{tr} \left((\alpha \alpha^T - K^{-1}) \frac{\partial K}{\partial \theta} \right), \alpha = K^{-1} y$$

Or we could also be bayesian about θ

4.4 Aproximation Techniques

Local method: $k(x_1, x_2) = 0$ if $\|x_1 - x_2\| \gg 1$

Random Fourier Features: if $k(x, y) = \kappa(x - y)$
 $p(w) = \mathcal{F} \{ \kappa(\cdot), w \}$. Then $p(w)$ can be normalized to be a density.

$$\kappa(x - y) = \mathbb{E}_{p(w)} \left[\exp \{ i w^T (x - y) \} \right] \text{ antitransform}$$

$$\kappa(x - y) = \mathbb{E}_{b \sim \mathcal{U}([0, 2\pi]), w \sim p(w)} [z_{w,b}(x) z_{w,b}(y)]$$

where $z_{w,b}(x) = \sqrt{2} \cos(w^T x + b)$. I can MC extract features z . If # features is $\ll n$ then this is faster ($X^T X$ vs XX^T)

Inducing points: We a vector of inducing variables u

$$f_A|u \sim N(K_{Au} K_{uu}^{-1} u, K_{AA} - K_{Au} K_{uu}^{-1} K_{uA})$$

$$f_*|u \sim N(K_{*u} K_{uu}^{-1} u, K_{**} - K_{*u} K_{uu}^{-1} K_{u*})$$

Subset of Regressors (SoR): $\blacksquare \rightarrow 0$

FITC: $\blacksquare \rightarrow$ its diagonal

5 Approximate inference

5.1 Laplace Approximation

$$\hat{\theta} = \arg \max_{\theta} p(\theta|y)$$

$$\Lambda = -\nabla_{\theta} \nabla_{\theta} \log p(\theta|y)|_{\theta=\hat{\theta}}$$

$$p(\theta|y) \simeq q(\theta) = N(\hat{\theta}, \Lambda^{-1})$$

5.2 Variational Inference

$$\hat{q} = \arg \min_{q \in \mathcal{Q}} KL(q||p(\cdot|y))$$

$$\hat{q} = \arg \max_{q \in \mathcal{Q}} ELBO \text{ Evidence Lower Bound}$$

$$ELBO \doteq \mathbb{E}_{\theta \sim q} [\log p(y|\theta)] - KL(q||p(\cdot)) \leq \log p(y)$$

5.3 Markov Chain Monte Carlo

Idea: All we need is sampling from posterior

Ergodic Markov Chain:

$$\exists t \text{ s.t. } \mathbb{P}(i \rightarrow j \text{ in } t \text{ steps}) > 0 \quad \forall i, j \Rightarrow$$

$$\exists ! \pi = \lim_{N \rightarrow \infty} \mathbb{P}(X_N = x) \text{ Limit distribution}$$

Ergodic Theorem: if $(X_i)_{i \in \mathbb{N}}$ is ergodic:

$$\lim_{N \rightarrow \infty} \frac{1}{N} \sum_{i=1}^N f(X_i) = \mathbb{E}_{x \sim \pi} [f(x)]$$

Detailed Blanced Equation:

$P(x|x')$ is the transition model of a MC:

if $R(x)P(x'|x) = R(x')P(x|x')$ then R is the limit distribution of the MC

Metropolis Hastings Algo: Sample from a MC

which has $P(x) = \frac{Q(x)}{Z}$ as limit dist.

Result: $\{X_i\}_{i \in \mathbb{N}}$ sampled from the MC

init: $R(x|x')$

/ Good R choice \rightarrow fast convergence */*

init: $X_0 = x_0$

for $t \leftarrow 1, 2, \dots$ **do**

$$x' \sim R(\cdot, x_{t-1})$$

$$\alpha = \min \left\{ 1, \frac{Q(x')R(x_{t-1}|x')}{Q(x_{t-1})R(x'|x_{t-1})} \right\}$$

with probability α **do**

$$\lfloor X_t = x';$$

$$\text{otherwise } X_t = x_{t-1};$$

Metropolis Adj. Langevin Algo (MALA):

$$\text{Energy function: } P(x) = \frac{Q(x)}{Z} = \frac{1}{Z} \exp(-f(x))$$

We chose: $R(x|x') = \mathcal{N}(x' - \tau \nabla f(x), 2\tau I)$

Stoch. Grad. Langevin Dynamics (SGLD):

We use SGD to Approximate ∇f . Converges also without acceptance step

Hamilton MC: SGD performance improved by adding momentum (consider last step ∇f)

Gibbs sampling: Practical when $X \in \mathbb{R}^n$

Used when $P(X_{1:n})$ is hard but $P(X_i|X_{-i})$ is easy.

init: $x_0 \in \mathbb{R}^n$; $(x_0^{(B)} = x^{(B)})$ B is our data

for $t = 1, 2, \dots$ **do**

with $i \sim \mathcal{U}(\{1 : n\} \setminus B)$ **do**

$$\lfloor x_{t-1}^{(i)} \sim P(x^{(i)}|x^{(-i)})$$

* if we do it $\forall i \notin B$ no DBE but more practical

5.4 Variable elimination for MPE (most probable explanation):

With loopy graphs, BP is often **overconfident/oscillates**.

6 Bayesian Neural Nets

Likelihood: $p(y|x; \theta) = \mathcal{N}(f_1(x, \theta), \exp(f_2(x, \theta)))$

Prior: $p(\theta) = \mathcal{N}(0, \sigma_p^2)$

$$\theta_{MAP} = \arg \max_{\theta} \log(p(y, \theta))$$

6.1 Variation inference:

Usually we use \mathcal{Q} = Set of Gaussians

$\hat{q} = \arg \max_{\mathcal{Q}} ELBO$ Reparameterization trick

q approx. the posterior but how to predict?

$$p(y^*|x^*, \mathcal{D}) \simeq \frac{1}{m} \sum_{j=1}^m p(y^*|x^*, \theta^{(j)}), \theta \sim \hat{q}(\theta)$$

Gaussian Mixture distribution: $\mathbb{V}(y^*|x^*, \mathcal{D}) \simeq$

$$\simeq \frac{1}{m} \sum_{j=1}^m \sigma^2(x^*, \theta^{(j)}) + \frac{1}{m} \sum_{j=1}^m (\mu(x^*, \theta^{(j)}) - \bar{\mu}(x^*))$$

■ \rightarrow Alestoric, ■ \rightarrow Epistemic

Dropouts Regularization: Random ignore nodes in SGD iteration: Equavalent to VI with

$$Q = \{q(\cdot|\lambda) = \prod_j q_j(\theta_j|\lambda), \lambda \in \mathbb{R}^d\}$$

where $q_j(\theta_j|\lambda) = p\delta_0(\theta_j) + (1-p)\delta_{\lambda_j}(\theta_j)$

This allows to do Dropouts also in prediction

6.2 MCMC:

MCMC but cannot store all the $\theta^{(i)}$:

1) Subsampling: Only store a subset of the $\theta^{(i)}$

2) Gaussian Aproximation: We only keep:

$$\mu_i = \frac{1}{T} \sum_{j=1}^T \theta_i^{(j)} \text{ and } \sigma_i = \frac{1}{T} \sum_{j=1}^T (\theta_i^{(j)} - \mu_i)^2$$

And update them online.

Predictive Esnable NNs:

Let $\mathcal{D} = \{(x_i, y_i)\}_{i=1:n}$ be our dataset.

Train θ_i^{MAP} on \mathcal{D}_i with $i = 1, \dots, m$

\mathcal{D}_i is a Bootstrap of \mathcal{D} of same size

and $p(y^*|x^*, \mathcal{D}) \simeq \frac{1}{m} \sum_{j=1}^m p(y^*|x^*, \theta_i^{MAP})$

6.3 Model calibration

Train \hat{q} on \mathcal{D}_{train}

Evaluate \hat{q} on $\mathcal{D}_{val} = \{(y', x')\}_{i=1:m}$

Held-Out-Likelihood $\doteq \log p(y'_{1:m}|x'_{1:m}, \mathcal{D}_{train})$

$$\geq \mathbb{E}_{\theta \sim \hat{q}} \left[\sum_{i=1}^m \log p(y'_i|x'_i, \theta) \right] \text{ (Jensen)}$$

$$\simeq \frac{1}{k} \sum_{j=1}^k \sum_{i=1}^m \log p(y'_i|x'_i, \theta^{(j)}), \theta^{(j)} \sim \hat{q}$$

Evaluate predicted accuracy: We divide \mathcal{D}_{val} into bins according to predicted confidence values. In each bin we compare accuracy with confidence

7 Active Learning

Let \mathcal{D} be the set of observable points.

We can observe $\mathcal{S} \subseteq \mathcal{D}, |\mathcal{S}| \leq R$

Information Gain: $\hat{\mathcal{S}} = \arg \max_{\mathcal{S}} F(\mathcal{S}) = I(f, y_{\mathcal{S}})$

For GPs: $F(\mathcal{S}) = \frac{1}{2} \log |I + \frac{1}{\sigma^2} K_{\mathcal{SS}}|$

This is NP Hard, \Rightarrow Greedy Algo:

init: $S^* = \emptyset$
for $t = 1 : R$ do
 $x_t = \arg \max_{x \in \mathcal{D}} F(S^* \cup \{x\})$
 $(x_t = \arg \max_{x \in \mathcal{D}} \sigma_x^2 |S \text{ for GPs})$
 $(x_t = \arg \max_{x \in \mathcal{D}} \frac{\sigma_{f|S}^2(x)}{\sigma_n^2(x)} \text{ for heter. GPs})$
 $S^* = S \cup \{x_t\}$

F is **Submodular** if: $\forall x \in \mathcal{D}, \forall A \subseteq B \subseteq \mathcal{D}$ holds that: $F(A \cup \{x\}) - F(A) \geq F(B \cup \{x\}) - F(B)$

F is Submodular $\Rightarrow F(S^*) \geq (1 - \frac{1}{e})F(\hat{S})$

8 Bayesian Optimization

Like Active Learning but we only want to find the optima. We pick x_1, x_2, \dots from \mathcal{D} and observe $y_i = f(x_i) + \epsilon_i$.

Cumulative regret: $R_T = \sum_{t=1}^T (\max_{x \in \mathcal{D}} f(x) - f(x_t))$

Oss: $\frac{R_T}{T} \rightarrow 0 \Rightarrow \max_t f(x_t) \rightarrow \max_{x \in \mathcal{D}} f(x)$

8.1 Upper Confidence Sampling

With GP $x_t = \arg \max_{x \in \mathcal{D}} \mu_{t-1}(x) + \beta_t \sigma_{t-1}(x)$

Choosing the correct β_t we get: $\frac{R_T}{T} = O\left(\sqrt{\frac{\gamma_T}{T}}\right)$.

Where $\gamma_t = \max_{|S| \leq T} I(f; y_S)$. On d dims:

Linear: $\gamma_T = O(d \log T)$ RBF: $\gamma_T = O((\log T)^{d+1})$

Optimal $\beta_t = O(\|f\|_K^2 + \gamma_t \log^3 T)$

Oss: $\beta \uparrow$ = more exploration

8.2 Thompson Sampling

$x_t = \arg \max_{x \in \mathcal{D}} \tilde{f}(x), \tilde{f} \sim p(f|x_{1:n}, y_{1:n})$

9 Markov Decision Process (MDP)

9.1 Definitions

$\mathcal{X} = \{1, \dots, n\}$ states; $\mathcal{A} = \{1, \dots, m\}$ actions;

$p(x'|x, a)$ transition probability;

$r(x, a)$ reward (can be random); $\pi: \mathcal{X} \rightarrow \mathcal{A}$ policy;

$T^\pi \in \mathbb{R}^{n \times n}, T_{ij}^\pi = p(j|i, \pi(i))$ Transition Matrix:

$J(\pi) = \mathbb{E}\left[\sum_{i=0}^{\infty} \gamma^i r(X_i, \pi(X_i))\right]$ Expected value:

$V^\pi: \mathcal{X} \rightarrow \mathbb{R}, x \mapsto J(\pi|X_0 = x)$ Value function;

$Q^V(x, a) = r(x, a) + \gamma \sum_{x' \in \mathcal{X}} p(x'|x, a) V(x')$ Q func;

$\pi_G^V(x) = \arg \max_a Q^V(x, a)$ greedy policy w.r.t. V ;

9.2 Value function Theorem

$V^\pi(x) = r(x, \pi(x)) + \gamma \sum_{x' \in \mathcal{X}} p(x'|x, \pi(x)) V^\pi(x')$

Matrix formulation: $(I - \gamma T^\pi) V^\pi = r^\pi$

9.3 Bellman Theorem

1) π^*, V^* are optimal policy and it's value func.

2) $\pi^* = \pi_G^{V^*}$

3) $V^*(x) = \max_a [r(x, a) + \gamma \sum_{x' \in \mathcal{X}} p(x'|x, a) V^*(x)]$

1) \Leftrightarrow 2) \Leftrightarrow 3)

9.4 Algorithms

9.4.1 Policy iteration

while no more changes **do**

$\pi \leftarrow \pi_G^V$ (Update the Policy)

$V \leftarrow (I - \gamma T^\pi)^{-1} r^\pi$ (Update the value)

9.4.2 Value iteration

while $\|V_t - V_{t-1}\| \leq \epsilon$ **do**

foreach $x \in \mathcal{X}, a \in \mathcal{A}$ **do**

$Q_t(x, a) \leftarrow r(x, a) + \gamma \sum_{x' \in \mathcal{X}} p(x'|x, a) V_{t-1}(x)$

foreach $x \in \mathcal{X}$ **do**

$V_t(x) \leftarrow \max_a Q_t(x, a)$

$\hat{\pi} = \pi_G^{V_T}$; where V_T last found Value

9.5 Partially Observable MDP (POMDP)

POMDP can be seen as MDP where:

1) \mathcal{X}_{POMDP} are prob. distribution over \mathcal{X}_{MDP}

2) the actions are the same

3) $r_{POMDP}(b, a) = \mathbb{E}_{x \sim b} [r_{MDP}(x, a)]$

4) Trans. model: $b_{t+1}(x) = \mathbb{P}(X_{t+1} = x | y_{1:t+1}, a_t)$

$b_{t+1}(x) = \frac{1}{Z} p(y_{t+1} | X_{t+1} = x) \sum_{x' \in \mathcal{X}_{MDP}} p(x|x', a_t) b_t(x')$

How to solve? Discretize \mathcal{X}_{POMDP} and treat it as a MDP or Policy gradient techniques

10 Non Parametric RL

It is an MDP with unknown $p(x'|x, a)$ and $r(x, a)$

10.1 Model-based RL

From all steps $X_{t+1}, R_t | X_t, A_t$ we can learn:

$p(x'|x, a) \approx \hat{p}_{x'|x, a} = \frac{\text{Count}(X_{t+1}=x', X_t=x, A_t=a)}{\text{Count}(X_t=x, A_t=a)}$

$r(x, a) \approx \hat{r}_{x, a} = \frac{1}{\text{Count}(X_t=x, A_t=a)} \sum_t |X_t=x, A_t=a R_t$

How to chose a_t ?

10.1.1 ϵ -greedy (On-Policy)

With probability ϵ , pick random action.

With probability $1 - \epsilon$, pick $a = \arg \max Q(x, a)$.

Oss: Q is calculated from (\hat{p}, \hat{r})

Th: If $\epsilon_t \xrightarrow{RM} 0$ then $(\hat{r}, \hat{p}) \xrightarrow{a.s.} (r, p)$

10.1.2 Softmax (On-Policy)

Draw $a \sim q(a|x) = \text{softmax} \frac{Q(x, a)}{\tau}$

If $\tau \uparrow$ it means I trust less Q

10.1.3 R_{max} algorithm (On-Policy)

We add a fairy state x^*

init: $r(x, a) = R_{max} \forall x \in \mathcal{X} \cup \{x^*\}, a \in \mathcal{A}$

init: $p(x^*|x, a) = 1 \forall x \in \mathcal{X}, a \in \mathcal{A}$

init: π = optimal policy w.r.t. p, r

repeat

 Execute π and get x_{t+1} and r_t

 Update belief of $r(x_t, \pi(x_t))$ and

$p(x_{t+1}|x_t, \pi(x_t))$

 If observed 'enough' in (x, a)

 recompute π using the updated belief only in (x, a)

until;

'Enough'? See Hoeffding's inequality

$(\hat{p} \in [0, 1], \hat{r} \in [0, R_{max}])$.

PAC bound: With probability $1 - \delta$, R_{max} will reach an ϵ -optimal policy in a number of steps that is polynomial in $|X|, |A|, T, 1/\epsilon$ and $\log(1/\delta)$.

Memory $O(|X|^2|A|)$.

10.2 Model-free RL

Learn π^* only via V^* or Q^{V^*}

10.2.1 TD-learning (On-Policy)

Given a policy π we want to learn V^π

$V^\pi(x) = \mathbb{E}_{R \sim r(x, \pi(x)), X' \sim p(\cdot|x, \pi(x))} [R + \gamma V^\pi(X')]$

After seeing $(x_{t+1}, r_t | x_t, \pi(x_t))$ we update:

$V_{t+1}(x_t) \leftarrow (1 - \alpha_t) V_t(x_t) + \alpha_t (r_t + \gamma V_t^\pi(x_{t+1}))$

Where α_t is a regularizer term (only 1 sample)

Th: If $\alpha_t \xrightarrow{RM} 0$ then $V \xrightarrow{a.s.} V^\pi$

10.2.2 Q-learning (Off Policy)

Given experience we want to learn $Q^* = Q^{V^*}$

$Q^*(x, a) = \mathbb{E}_{R \sim r(x, \pi(x)), X' \sim p(\cdot|x, \pi(x))} [R + \gamma \max_{a'} Q^*(X', a')]$

After seeing $(x_{t+1}, r_t | x_t, a_t)$ we update:

$Q(x_t, a_t) \leftarrow (1 - \alpha_t) Q(x_t, a_t) + \alpha_t (r_t + \gamma \max_{a'} Q(x_{t+1}, a'))$

Th: If $\alpha_t \xrightarrow{RM} 0$ then $Q \xrightarrow{a.s.} Q^*$

Optimistic Q learning:

Initialize: $Q(x, a) = \frac{R_{max}}{1 - \gamma} \prod_{t=1}^{T_{init}} (1 - \alpha_t)^{-1}$

Same convergence time as with R_{max} . Memory $O(|X||A|)$. Comp: $O(|A|)$.

11 Parametric RL

11.1 Parametric TD-learning

11.1.1 TD-learning as SGD

TD-learning = 1 sample $(x', r | x, \pi(x))$ SGD on:

$\bar{l}_2(V; x, r) = \frac{1}{2} (V - r - \gamma \mathbb{E}_{x' \sim p(\cdot|x, \pi(x))} [\hat{V}^\pi(x')])^2$

1 sample estimate of $\nabla_V \bar{l}_2 = \delta = V - r - \gamma \hat{V}^\pi(x')$

$\Rightarrow V \leftarrow V - \alpha_t \delta$ where $V = \hat{V}^\pi(x)$

11.1.2 TD-parametric

If $\hat{V}^\pi(x) = V(x, \theta)$ then:

$\delta = [V(x; \theta) - r - \gamma V(x'; \theta_{old})] \nabla_\theta V(x, \theta)$

11.2 Parametric Q-learning

$\delta(\theta, \theta_{old}) = (Q(x, a; \theta) - r - \gamma \max_{a'} Q(x', a'; \theta_{old}))$

We don't differentiate with regard to θ_{old}

The SGD step is: $\theta \leftarrow \theta - \alpha_t \delta(\theta, \theta) \nabla_\theta Q(x, a; \theta)$

Deep Q Networks (DQN): Version of Q-learning where we update Q only each batch:

$L(\theta) = \sum_{(x, a, r, x') \in \mathcal{D}} (r + \gamma \max_{a'} Q(x', a'; \theta_{old}) - Q(x, a; \theta))^2$

Double DQN (better):

$L(\theta) = \sum_{(x, a, r, x') \in \mathcal{D}} (r + \gamma Q(x', a^*(\theta); \theta_{old}) - Q(x, a; \theta))^2$

where: $a^*(\theta) \doteq \arg \max_{a'} Q(x', a'; \theta)$

11.3 Policy-Search method