

**1 Bayesian Regression**  
 $w \sim N(0, \sigma_p^2 I)$ ,  $\epsilon \sim N(0, \sigma_n^2 I)$ ,  $y = Xw + \epsilon$   
 $y|w \sim N(Xw, \sigma_n^2 I)$   
 $w|y \sim N((X^T X + \lambda I)^{-1} X^T y, (X^T X + \lambda I)^{-1} \sigma_n^2)$

**2 Kalman Filter**  
 $\begin{cases} X_{t+1} = F X_t + \epsilon_t & \epsilon_t \sim N(0, \Sigma_x) \\ Y_t = H X_t + \eta_t & \eta_t \sim N(0, \Sigma_y) \end{cases}$   $X_1 \sim N(\mu_p, \Sigma_p)$   
Then if  $X_0$  is Gaussian then  $X_t|Y_{1:t} \sim N(\mu_t, \Sigma_t)$ :  
 $\mu_{t+1} = F \mu_t + K_{t+1}(y_{t+1} - H F \mu_t)$   
 $\Sigma_{t+1} = (I - K_{t+1} H)(F \Sigma_t F^T + \Sigma_x)$   
 $K_{t+1} = (F \Sigma_t F^T + \Sigma_x) H^T (H(F \Sigma_t F^T + \Sigma_x) H^T + \Sigma_y)^{-1}$

**3 Gaussian Processes**  
 $f \sim GP(\mu, k) \Rightarrow \forall \{x_1, \dots, x_n\} \forall n < \infty$   
 $[f(x_1) \dots f(x_n)] \sim N([\mu(x_1) \dots \mu(x_n)], K)$   
where  $K_{ij} = k(x_i, x_j)$

**3.1 Gaussian Process Regression**  
 $f \sim GP(\mu, k)$  then:  $f|y_{1:n}, x_{1:n} \sim GP(\tilde{\mu}, \tilde{k})$   
 $\tilde{\mu}(x) = \mu(x) + K_{A,x}^T (K_{AA} + \epsilon I_n)^{-1} (y_A - \mu_A)$   
 $\tilde{k}(x, x') = k(x, x') - K_{A,x}^T (K_{AA} + \epsilon I_n)^{-1} K_{A,x'}$   
Where:  $K_{A,x} = [k(x_1, x) \dots k(x_n, x)]^T$   
 $[K_{AA}]_{ij} = k(x_i, x_j)$  and  $\mu_A = [\mu(x_1) \dots \mu(x_n)]^T$

**3.2 Kernels**  
 $k(x, y)$  is a kernel if it's symmetric semidefinite positive:  
 $\forall \{x_1, \dots, x_n\}$  then for the Gram Matrix  
 $[K]_{ij} = k(x_i, x_j)$  holds  $c^T K c \geq 0 \forall c$   
**Some Kernels:** (h is the bandwidth hyperp.)  
Gaussian (rbf):  $k(x, y) = \exp(-\frac{\|x-y\|^2}{h^2})$   
Exponential:  $k(x, y) = \exp(-\frac{\|x-y\|}{h})$   
Linear kernel:  $k(x, y) = x^T y$  (here  $K_{AA} = X X^T$ )

**3.3 Optimization of Kernel Parameters**  
Given a dataset  $A$ , a kernel function  $k(x, y; \theta)$ .  
 $y \sim N(0, K_y(\theta))$  where  $K_y(\theta) = K_{AA}(\theta) + \sigma_n^2 I$   
 $\hat{\theta} = \arg \max_{\theta} \log p(y|X; \theta)$   
In GP:  $\hat{\theta} = \arg \min_{\theta} y^T K_y^{-1}(\theta) y + \log |K_y(\theta)|$   
We can from here  $\nabla \downarrow$ :  
 $\nabla_{\theta} \log p(y|X; \theta) = \frac{1}{2} \text{tr}((\alpha \alpha^T - K^{-1}) \frac{\partial K}{\partial \theta})$ ,  $\alpha = K^{-1} y$   
Or we could also be bayesian about  $\theta$

**3.4 Aproximation Techniques**  
**Local method:**  $k(x_1, x_2) = 0$  if  $\|x_1 - x_2\| \gg 1$   
**Random Fourier Features:** if  $k(x, y) = \kappa(x - y)$   
 $p(w) = \mathcal{F} \{\kappa(\cdot), w\}$ . Then  $p(w)$  can be normalized to be a density.  
 $\kappa(x - y) = \mathbb{E}_{p(w)} [\exp\{i w^T (x - y)\}]$  antitransform  
 $\kappa(x - y) = \mathbb{E}_{b \sim \mathcal{U}([0, 2\pi]), w \sim p(w)} [z_{w,b}(x) z_{w,b}^*(y)]$

where  $z_{w,b}(x) = \sqrt{2} \cos(w^T x + b)$ . I can MC extract features  $z$ . If # features is  $\ll n$  then this is faster ( $X^T X$  vs  $X X^T$ )

**Inducing points:** We a vector of inducing variables  $u$   
 $f_A|u \sim N(K_{Au} K_{uu}^{-1} u, K_{AA} - K_{Au} K_{uu}^{-1} K_{uA})$   
 $f_*|u \sim N(K_{*u} K_{uu}^{-1} u, K_{**} - K_{*u} K_{uu}^{-1} K_{u*})$

**Subset of Regressors (SoR):**  $\blacksquare \rightarrow 0$   
**FITC:**  $\blacksquare \rightarrow$  its diagonal

**4 Review of useful concepts and Introduction**

**4.1 Multivariate Gaussian**  
 $f(x) = \frac{1}{2\pi\sqrt{|\Sigma|}} e^{-\frac{1}{2}(x-\mu)^T \Sigma^{-1} (x-\mu)}$   
Suppose we have a Gaussian random vector  $X_V \sim N(\mu_V, \Sigma_{VV})$ .  
Suppose we take two disjoint subsets of  $V$ :  $A = i_1, \dots, i_k$  and  $B = j_1, \dots, j_m$ .  
Then, the conditional distribution:  $P(X_A|X_B = x_B) = N(\mu_{A|B}, \Sigma_{A|B})$  is Gaussian:  
 $\mu_{A|B} = \mu_A + \Sigma_{AB} \Sigma_{BB}^{-1} (x_B - \mu_B)$   
 $\Sigma_{A|B} = \Sigma_{AA} - \Sigma_{AB} \Sigma_{BB}^{-1} \Sigma_{BA}$

**4.2 Convex / Jensen's inequality**  
 $g(x)$  is convex  $\Leftrightarrow x_1, x_2 \in \mathbb{R}, \lambda \in [0, 1] : g''(x) > 0$   
 $g(\lambda x_1 + (1 - \lambda)x_2) \leq \lambda g(x_1) + (1 - \lambda)g(x_2)$   
 $\varphi(\mathbb{E}[X]) \leq \mathbb{E}[\varphi(X)]$

**4.3 Review Probability**  
**Probability space**  $(\Omega, F, P)$ : Set of atomic events  $\Omega$ . Set of all non-atomic events ( $\sigma$ -Algebra):  $F \in 2^\Omega$ . Probability measure:  $P : F \rightarrow [0, 1]$   
**Bayes' rule:**  $P(B|A) = \frac{P(A, B)}{P(A)} = \frac{P(A|B)P(B)}{P(A)}$ , where  $P(A) = \sum_b P(A|B)P(B)$   
**Union:**  $P(A \cup B) = P(A) + P(B) - P(A \cap B)$   
**Rules for joint distributions:**  
Sum rule (Marginalization):  $P(X_{1:i-1}, X_{i+1:n}) = \sum_{x_i} P(X_{1:i-1}, X_i = x_i, X_{i+1:n})$   
Product rule (Chain rule):  $P(X_{1:n}) = P(x_1)P(x_2|x_1) \dots P(x_n|x_{1:n-1})$   
**Conditional Independence:**  
 $X \perp Y|Z$  iff  $P(X, Y|Z) = P(X|Z)P(Y|Z)$   
If  $P(Y|Z) > 0 \Rightarrow P(X|Z, Y) = P(X|Z)$   
**Properties of Conditional Independence:**  
Symmetry:  $X \perp Y | Z \Rightarrow Y \perp X | Z$   
Decomposition:  $X \perp (Y, W) | Z \Rightarrow X \perp Y | Z$   
Contraction:  $(X \perp Y | Z) \wedge (X \perp W | Y, Z) \Rightarrow X \perp Y, W | Z$   
Weak union:  $X \perp Y, W | Z \Rightarrow X \perp Y | Z, W$   
Intersection:  $(X \perp Y | W, Z) \wedge (X \perp W | Y, Z) \Rightarrow X \perp Y, W | Z$

**5 Bayesian Networks**

**5.1 Basic concepts**  
A Bayesian network  $(G, P)$  consists of:  
- A BN structure  $G$  (directed, acyclic graph)

- Set of conditional probability distributions  
 $(G, P)$  defines the joint distribution:  
 $P(X_1, \dots, X_n) = \prod_i P(X_i | Pa_{X_i})$   
BNs with 3 nodes:  
**5.2 Active trails and d-separation**  
An undirected path in a BN structure  $G$  is called active trail for observed variables  $O \in X_1, \dots, X_n$  if for every consecutive triple of variables  $X, Y, Z$  on the path:  
- **indirect causal effect:**  $X \rightarrow Y \rightarrow Z$  and  $Y$  unobserved  
- **indirect evidential effect:**  $X \leftarrow Y \leftarrow Z$  and  $Y$  unobserved  
- **common cause:**  $X \leftarrow Y \rightarrow Z$  and  $Y$  unobserved.  
- **common effect:**  $X \rightarrow Y \leftarrow Z$  and  $Y$  or any of  $Y$ 's descendants is observed.  
Any variables  $X_i$  and  $X_j$  for which there is no active trail for observations  $O$  are called d-separated by  $O$ .  
**Theorem:**  $d\text{-sep}(X_i; X_j | O) \Rightarrow X \perp Y | Z$   
Converse does not hold in general!

**6 Exact inference (tree-structured BN)**

**6.1 Variable elimination**  
- Given a BN and query  $P(X|E = e)$   
- Choose an ordering of  $X_1, \dots, X_n$  **Eliminate variables from the outside in!**  
- Set up initial factors:  $f_i = P(X_i | Pa_i)$   
- For  $i = 1 : n, X_i \notin X, E$   
- Collect and multiply all factors  $f$  that include  $X_i$   
- Generate new factor by marginalizing out  $X_i$ :  $g_{X_i} = \sum_{x_i} \prod_j f_j$   
- Add  $g$  to set of factors  
- Renormalize  $P(x, e)$  to get  $P(x|e)$   
**Variable elimination for polytrees:**  
- Pick a root, (avoiding  $X$  and  $E$ )  
- Orient edges towards root  
- Eliminate variables according to topological order

**6.2 Avoiding recomputation: factor graphs**  
FG for a BN is a bipartite graph consisting of variables (circles) and factors (rectangles). **It is not a unique representation.**

**6.2.1 Sum-product/Belief Propagation (BP) Algorithm:**  
- Initialize all messages as uniform distribution  
- Until converged to:  
- Pick a root in the factor graph and reorient the edges towards this root.

- Update messages according to this ordering. Do passes from leaves to root and from root to leaves.  
- If a leaf node is a variable node:  $\mu_{x \rightarrow f}(x) = 1$   
- If a leaf node is a factor node:  $\mu_{f \rightarrow x}(x) = f(x)$   
Messages from node  $v$  to factor  $u$ :  
 $\mu_{v \rightarrow u}(x_v) = \prod_{u' \in N(v) \setminus \{u\}} \mu_{u' \rightarrow v}(x_v)$   
- Messages from factor  $u$  to node  $v$ :  
 $\mu_{u \rightarrow v}(x_v) = \sum_{x_{u'} \sim x_v} f_u(x_u) \prod_{v' \in N(u) \setminus \{v\}} \mu_{v' \rightarrow u}(x_{v'})$   
- Break once all messages change by  $\leq \epsilon$   
**Hope:** after convergence, we have:  
 $P(X_v = x_v) = \frac{1}{Z} \prod_{u \in N(v)} \mu_{u \rightarrow v}(x_v)$   
 $P(\bar{X}_u = \bar{x}_u) = \frac{1}{Z} f_u(\bar{x}_u) \prod_{v \in N(u)} \mu_{v \rightarrow u}(x_v)$   
**If we have a polytree Bayesian network:**  
- Choose one node as root  
- Send messages from leaves to root and from root to leaves

**7 Approximate inference (loopy networks)**  
With loopy graphs, BP is often **overconfident/oscillates**.

**7.1 Variable elimination for MPE (most probable explanation):**  
- Given BN and evidence  $E=e$   
- Choose an ordering of  $X_1, \dots, X_n$   
- Set up initial factors  $f_i = P(X_i | Pa_i)$   
- For  $i = 1 : n, X_i \notin E$ :  
- Collect and multiply all factors  $f_j$  that include  $X_i$   
- Generate new factor by maximizing out  $X_i$ :  $g_i = \max_{w=x_i} \prod_j f_j$   
- Add  $g$  to set of factors  
- For  $i = n-1 : 1, X_i \notin E$ :  $\hat{x}_i = \arg \max_{x_i} g_i(x_i, \hat{x}_{i+1:n})$

**Retrieving MAP from Max-Product (MAP = MPE for a subset of RVs):**  
- Define max-marginals:  
 $P_{\max}(X_v = x_v) := \max_{x \sim x_v} P(x)$   
- For tree factor graphs, max-product computes max-marginals:  
 $P_{\max}(X_v = x_v) \propto \prod_{u \in N(v)} \mu_{u \rightarrow v}(x_v)$   
- Can retrieve MAP solution from these (must be careful when ties need to be broken).

**7.2 Sampling based inference: compute marginals as expectations**  
**Hoeffding's inequality:** Suppose  $f$  is bounded in  $[0, C]$ . Then:  
 $P(|\mathbb{E}_p[f(X)] - \frac{1}{N} \sum_{i=1}^N f(x_i)| \geq \epsilon) \leq 2 \exp(-\frac{2N\epsilon^2}{C^2})$   
**Monte Carlo Sampling from a BN:**  
- Sort variables in topological ordering  $X_1, X_n$   
- For  $i = 1$  to  $n$ , sample:  
 $x_i \sim P(X_i = x_i | x_{1:i-1})$

## Bayesian Sampling:

- Collect samples over all variables:

$$\hat{P}(X_A = x + A | X_B = x_B) = \frac{\text{Count}(x_A, x_B)}{\text{Count}(x_B)}$$

- Throw away samples that disagree with  $x_B$
- Count fraction of  $x_a$  on remaining samples

### 7.2.1 Directly sampling from the posterior: MCMC

**Markov Chain::** A (stationary) MC is a sequence of RVs  $X_1, \dots, X_N$ , with prior  $P(X_1)$  and transition probabilities  $P(X_{t+1}|X_t)$  independent of  $t$ .

**Markov assumpt.:**  $X_{1:t-1} \perp X_{t+1:T} | X_t, \forall t > 1$   
**Stationarity assumption:**

$$P(X_{t+1} = x | X_t = x') = P(X_t = x | X_{t-1} = x'), \forall t > 1$$

**If ergodic** (= there exists a finite  $t$  such that every state can be reached in exactly  $t$  steps), then: it has a unique and positive stationary distribution  $\pi(X) > 0$ , such for all  $x$ :

$$\lim_{N \rightarrow \infty} P(X_N = x) = \pi(x) \text{ and } \pi(X) \perp P(X_1).$$

If MC satisfies the **detailed balance equation** (for unnormalized distribution  $Q$ , for all  $x, x'$ :  $Q(x)P(x'|x) = Q(x')P(x|x')$ ), then the MC has stationarity distribution  $\pi(X) = 1/ZQ(X)$ .

### Designing Markov Chains:

- Proposal distribution  $R(X'|X)$ : given  $X_t = x$ , sample "proposal"  $x' \sim R(X'|X = x)$
- Acceptance distribution
  - Suppose  $X_t = x$
  - With probability  $\alpha = \min\left\{1, \frac{Q(x')R(x|x')}{Q(x)R(x'|x)}\right\}$

- set:  $X_{t+1} = x'$
- With probability  $1 - \alpha$ , set  $X_{t+1} = x$

### MCMC for graphical models: Gibbs sampling (Random Vs **Practical variant**):

- Start with initial assignment  $x$  to all variables
- Fix observed variables  $X_B = x_B$
- For  $t = 1$  to  $\infty$ , do:
  - Pick a variable  $i$  uniformly at random from  $\{1, \dots, n\} \setminus B$  / **Set ordering**, and then, for each  $X_i$  (except those in  $B$ )
  - Set  $v_i =$  values of all  $x$  except  $x_i$
  - Sample  $x_i$  from  $P(X_i | v_i)$

### 8 Dynamical models (include time)

**8.1 Examples with one variable per time step**  
 $X_1, \dots, X_T$  (unobserved) hidden states  
 $Y_1, \dots, Y_T$  (noisy) observations

**HMMs (polytrees: can use belief propagation):**  $X_i$  categorical,  $Y_i$  categorical (or arbitrary)  
**Kalman filters:**  $X_i, Y_i$  Gaussian distributions  
-  $P(X_1)$ : prior belief about location at time  $i$   
-  $P(X_{t+1}|X_t)$ : **'Motion model'** (how do I expect my target to move in the environment?):  
 $X_{t+1} = FX + \epsilon_t$  where  $\epsilon_t \sim N(0, \Sigma_x)$

-  $P(Y_t|X_t)$ : **'Sensor model'** (what do I observe if target is at location  $X_t$ ?)  $Y_t = HX_t + \eta_t$  where  $\eta_t \sim N(0, \Sigma_y)$

### 8.2 Inference tasks

**Filtering:**  $P(X_t | y_{1:t})$  Is it raining today?

**Prediction:**  $P(X_{t+\tau} | Y_{1:t})$  Rain 5 days from now?  
Example for one step:  $P(X_{t+1} | Y_{1:t}) = \sum_x P(X_{t+1}, X_t = x | Y_{1:t}) = \sum_x P(X_{t+1} | X_t = x_t) P(X_t | Y_{1:t})$  (with KFs, you need **integrals**!)

**Smoothing:**  $P(X_t | y_{1:t})$  with  $\tau < t$  Did it rain last week? [Can use sum-product (aka forward-backward).]

**MPE:**  $\underset{x_{1:T}}{\operatorname{argmax}} P(x_{1:T} | y_{1:T})$  Can use max product (aka Viterbi algorithm).

**Bayesian filtering:** Start with  $P(X_1)$ :

At time  $t$ , assume we have  $P(X_t | y_{1:t-1})$

$$\text{Conditioning: } P(X_t | y_{1:t}) = \frac{P(X_t | y_{1:t-1}) P(y_t | X_t)}{\sum_{x_t} P(X_t | y_{1:t-1}) P(y_t | X_t)}$$

Prediction ( $O(n^2)$  vs  $O(n)$  in conditioning):

$$P(X_{t+1} | y_{1:t}) = \sum_x P(X_{t+1} | X_t = x) P(X_t | y_{1:t})$$

**Since HMM is a polytree, smoothing/MPE can be computed by VE/BP. Kalman filtering:** Bayesian filtering for continuous problems. RV corrupted by Gaussian distributions with zero mean. **Bayesian filtering is basically the same, except that sums turn to integrals. General Kalman update**

- Transition model:  $P(x_{t+1} | x_t) = N(x_{t+1}; Fx_t, \Sigma_x)$
- Sensor model:  $P(y_t | x_t) = N(y_t; Hx_t, \Sigma_y)$

- Kalman update:

$$\mu_{t+1} = F\mu_t + K_{t+1}(y_{t+1} - HF\mu_t)$$

$$\Sigma_{t+1} = (I - K_{t+1})(F\Sigma_t F^T + \Sigma_x)$$

$$\text{- Kalman gain: } K_{t+1} = (F\Sigma_t F^T + \Sigma_x) H^T (H(F\Sigma_t F^T + \Sigma_x) H^T + \Sigma_y)^{-1}$$

### 8.3 Examples with > 1 variable per time step

**Dynamic Bayesian Networks:** a BN at every time step

These models typically have many loops. Exact inference is usually intractable.

### 8.4 Approx. infer. for filtering (DBNs and non-linear Kalman filters): Particle filtering

**Suppose:**  $P(X_t | y_{1:t}) \approx \frac{1}{N} \sum_{i=1}^N \delta_{x_{i,t}}$ , where  $\delta$  is the indicator function. **Prediction:** Propagate each particle:  $x'_i \sim P(X_{t+1} | x_{i,t})$

**Conditioning:**

- weight particles  $w_i = \frac{1}{Z} P(y_{t+1} | x'_i)$
- resample  $N$  particles  $x_{i,t+1} \sim \frac{1}{Z} \sum_{i=1}^N w_i \delta_{x'_i}$

**Conclusion we came to:**  $Z = \sum_{i=1}^N w_i \delta_{x_i}$

## 9 Probabilistic Planning

### 9.1 Markov Decision Processes

An MDP is specified by a quintuple:  $(X, A, r, P(x'|x, a), \gamma)$ , where  $X$  are states,  $A$  are actions,  $r(x, a)$  is a reward function and transition probabilities:

$$P(x'|x, a) = \text{Prob}(\text{Next state} = x' | \text{Action } a)$$

**Objective:** find a stationary policy  $\pi : S \rightarrow A$  that maximizes the sum of cumulative rewards.

**Value of a state given a policy:** sum of cumulative rewards, given that the initial state is this state  $\rightarrow$  **Bellman equation:**

$$V^\pi(s) = \mathbb{E} \left[ \sum_{t=0}^{\infty} \gamma^t r(s_t, \pi(s_t), s_{t+1}) | s_0 = s \right]$$

$$= \sum_{s' \in S} P(s' | s, \pi(s)) [r(s, \pi(s), s') + \gamma V^\pi(s')]$$

$$= r(s, \pi(s)) + \gamma \sum_{s' \in S} P(s' | s, \pi(s)) V^\pi(s')$$

**Theorem (Bellman):** a policy is optimal iff it is greedy w.r.t. its induced value function!

$$V^*(x) = \max_a [r(x, a) + \gamma \sum_{x'} P(x' | x, a) V^*(x')]$$

Bellman equation mais geral:

$$V^*(x) = \max_a [\sum_{x'} P(x' | x, a) (r(a, x, x') + \gamma V^*(x'))]$$

Optimal policy:

$$\pi^*(s) = \underset{a \in A}{\operatorname{argmax}} [r(x, a) + \gamma \sum_{x'} P(x' | x, a) V^*(x')]$$

### 9.2 Policy iteration (Cost $O(S^3 + SA\Delta)$ )

Start with an arbitrary (e.g. random) policy  $\pi$ . Until converged, do:

- Compute value function  $V^\pi(x)$
- Compute greedy policy  $\pi_G$  w.r.t.  $V^\pi$
- Set  $\pi \leftarrow \pi_G$

Guaranteed to monotonically improve and to converge to an **optimal** policy  $\pi^*$  in  $O(n^2 m / (1 - \gamma))$  iterations (converges in polynomial number of iterations)!

### 9.3 Value iteration (Cost $O(SA\Delta)$ )

Initialize  $V_0(x) = \max_a r(x, a)$

For  $t = 1$  to  $\infty$ :

- For each  $(x, a)$ , let:

$$Q_t(x, a) = r(x, a) + \gamma \sum_{x'} P(x' | x, a) V_{t-1}(x')$$

- For each  $x$ , let  $V_t(x) = \max_a Q_t(x, a)$

- Break if  $\|V_t - V_{t-1}\|_\infty = \max_x |V_t(x) - V_{t-1}(x)| \leq \epsilon$

Then choose greedy policy w.r.t.  $V_t$ .

Guaranteed to converge to  $\epsilon$ -optimal policy (finds approximate solution in polynomial number of iterations)!

### 9.4 POMDP = Belief-state MDP

States = beliefs over states for original POMDP  
 $B = \Delta(1, \dots, n) = \{b : 1, \dots, n \rightarrow [0, 1], \sum_x b(x) = 1\}$

Actions: same as original MDP

**Transition model:**

- Stochastic observation:

$$P(Y_t | b_t) = \sum_{x=1}^n P(Y_t | X_t = x) b_t(x)$$

- State update (Bayesian filtering!): given

$$b_t, y_t, a_t: b_{t+1}(x') = \frac{1}{Z} \sum_x b_t(x) P(y_t | x) P(X_{t+1} =$$

$$x' | X_t = x, a_t)$$

Reward function:  $r(b_t, a_t) = \sum_x b_t(x) r(x, a_t)$

### 9.5 Example of approx. solution to POMDPs: Policy gradients

- Assume parameterized policy:  $\pi(b) = \pi(b; \theta)$
- For each parameter  $\theta$  the policy induces a Markov chain
- Can compute expected reward  $J(\theta)$  by sampling.
- Find optimal parameters through search (gradient ascent):  $\theta^* = \underset{\theta}{\operatorname{argmax}} J(\theta)$

## 10 Learning models from training data

### 10.1 Learning from i.i.d data

**Algorithm for Bayes Net MLE:**

Given BN of structure  $G$  and dataset  $D$  of complete observations

$$\text{For each } X_i \text{ estimate: } \hat{\theta}_{X_i | Pa_i} = \frac{\text{Count}(X_i, Pa_i)}{\text{Count}(Pa_i)}$$

Pseudo-counts for lime and cherry flavor:

$$\theta_{F=c} = \frac{\text{Count}(F=c) + \alpha_c}{N + \alpha_c + \alpha_l}$$

#### 10.1.1 Score based structure learning

Define scoring function  $S(G; D)$  and search over BN structure  $G$ :  $G^* = \underset{G}{\operatorname{argmax}} S(G; D)$

**Examples of scores:**

**MLE Score:**

$$\log P(D | \theta_G, G) = N \sum_{i=1}^n \hat{I}(X_i; Pa_i) + \text{const.}$$

**Where mutual information** ( $I(X_i, X_j) \geq 0$ ) is:

$$I(X_i, X_j) = \sum_{x_i, x_j} P(x_i, x_j) \log \frac{P(x_i, x_j)}{P(x_i)P(x_j)}$$

**Empirical mutual information:**

$$\hat{P}(x_i, x_j) = \frac{\text{Count}(x_i, x_j)}{N}$$

$$\hat{I}(X_i, X_j) = \sum_{x_i, x_j} \hat{P}(x_i, x_j) \log \frac{\hat{P}(x_i, x_j)}{\hat{P}(x_i)\hat{P}(x_j)}$$

**Regularizing a Bayes Net:**

$$S_{BIC}(G) = \sum_{i=1}^n \hat{I}(X_i; Pa_i) - \frac{\log N}{2N} |G|$$

where  $G$  is the number of parameters,  $n$  the number of variables and  $N$  the number of training examples.

**Chow-Liu algorithm:**

- For each pair  $X_i, X_j$  of variables, compute:  $\hat{P}(x_i, x_j) = \frac{\text{Count}(x_i, x_j)}{N}$
- Compute mutual information
- Define complete graph with weight of edge  $(X_i, X_j)$  given by the mutual information
- Find max spanning tree  $\rightarrow$  undirected tree
- Pick any variable as root and orient the edges away using breadth-first search.

## 11 Reinforcement Learning

### 11.1 Model-based RL

**11.1.1  $\epsilon$  greedy**  
 With probability  $\epsilon$ , pick random action. With prob  $(1 - \epsilon)$ , pick best action. If sequence  $\epsilon$  satisfies Robbins Monro criteria  $\rightarrow$  convergence to optimal policy with prob 1.

### 11.1.2 $R_{max}$ algorithm

**Input:** starting  $x_0$ , discount factor  $\gamma$ .

**Initially:** add fairy tale state  $x^*$  to MDP

- Set  $r(x, a) = R_{max}$  for all states  $x$  and actions  $a$

- Set  $P(x^* | x, a) = 1$  for all states  $x$  and actions  $a$

- Choose the optimal policy for  $r$  and  $P$

**Repeat:** 1. Execute policy  $\pi$  and, for each visited state/action pair, update  $r(x, a)$

2. Estimate transition probabilities  $P(x' | x, a)$

3. If observed 'enough' transitions/rewards, recompute policy  $\pi$ , according to current model  $P$  and  $r$ .

**Enough?** See Hoeffding's inequality. To reduce error  $\epsilon$ , need more samples  $N$ .

**Theorem:** With probability  $1 - \delta$ ,  $R_{max}$  will reach an  $\epsilon$ -optimal policy in a number of steps that is polynomial in  $|X|, |A|, T, 1/\epsilon$  and  $\log(1/\delta)$ .

Memory  $O(|X|^2 |A|)$ .

## 11.2 Model-free RL: estimate $V^*(x)$ directly

### 11.2.1 Q-learning

$Q(x, a) \leftarrow (1 - \alpha_t)Q(x, a) + \alpha_t(r + \gamma \max_{a'} Q(x', a'))$

**Theorem:** If learning rate  $\alpha_t$  satisfies:

$\sum_t \alpha_t = \infty$  and  $\sum_t \alpha_t^2 < \infty$  (Robbins-Monro), and actions are chosen at random, then  $Q$  learning converges to optimal  $Q^*$  with probability 1.

**Optimistic Q learning:**

Initialize:  $Q(x, a) = \frac{R_{max}}{1 - \gamma} \prod_{t=1}^{T_{init}} (1 - \alpha_t)^{-1}$

Same convergence time as with  $R_{max}$ . Memory  $O(|X||A|)$ . Comp:  $O(|A|)$ .

**Parametric Q-function approximation:**

$Q(x, a; \theta) = \theta^T \phi(x, a)$  to scale to large state spaces. (You can use Deep NN here!)

**SGD for ANNs:** initialize weights. For  $t = 1, 2, \dots$ , pick a data point  $(x, y)$  uniformly at random. Take step in negative gradient direction. (In practise, mini-batches).

**Deep Q Networks:** use CNN to approx Q function.  $L(\theta) = \sum_{(x, a, r, x') \in D} (r + \gamma \max_{a'} Q(x', a'; \theta^{old}) - Q(x, a; \theta))^2$

**Double DQN:** current network for evaluating argmax (too optimistic, and you remove  $\theta^{old}$  and put  $\theta$ ).

## 11.3 Gaussian processes

A GP is an (infinite) set of random variables (RV), indexed by some set  $X$ , i.e., for each  $x$  in  $X$ , there is a RV  $Y_x$  where there exists functions  $\mu : X \rightarrow \mathbb{R}$  and  $K : X \times X \rightarrow \mathbb{R}$  such that for all:  $A \in X$ ,  $A = x_1, \dots, x_k$ , it holds that

$Y_A = [Y_{x_1}, \dots, Y_{x_k}] \sim N(\mu_A, \Sigma_{AA})$ , where:  $\Sigma_{AA} =$  matrix with all combinations of  $K(x_i, x_j)$ .

$K$  is called kernel (covariance) function (must be symmetric and pd) and  $\mu$  is called mean function. **Making prediction with GPs:**

Suppose  $P(f) = GP(f; \mu, K)$  and we observe  $y_i = f(\vec{x}_i) + \epsilon_i$ ,  $A = \{\vec{x}_1 : \vec{x}_k\}$   $P(f(x) | \vec{x}_1 : \vec{x}_k, y_{1:k}) = GP(f; \mu', K')$ . In particular,  $P(f(x) | \vec{x}_1 : \vec{x}_k, y_{1:k}) = N(f(x); \mu_{x|A}, \sigma_{x|A}^2)$ ,

where  $\mu_{x|A} = \mu(\vec{x}) + \Sigma_{x,A}(\Sigma_{AA} + \sigma^2 I)^{-1} \Sigma_{x,A}^T (\vec{y}_A - \mu_A)$  and  $\sigma_{x|A}^2 = K(\vec{x}, \vec{x}) - \Sigma_{x,A}(\Sigma_{AA} + \sigma^2 I)^{-1} \Sigma_{x,A}^T$ .

**Closed form formulas for prediction!**